

## Nolan Anderson | Week 2 Assignment 1 | CPE 613 | January 21 2023

- Implement the RGB to grayscale kernel and run it on a test JPEG image. Use MATLAB or Python functions to generate the RGB data from an image and view the before and after results. Upload a brief PDF report describing your process and showing your results.
- Use this image: [goats.jpeg](#)
- In Matlab, load it via `rgb = imread('goats.jpeg')` to get the RGB in a 456 x 810 x 3 uint8.
- Run `imshow(rgb)` to convince yourself this is an image of goats.
- Write this out to a file and load it into your program.
- Copy it up to the device.
- Convert to grayscale on the device.
- Copy the result down.
- Write the result to a file.
- Load the file into Matlab.
- Run `imshow` on the loaded data to see the grayscale image.

### Before and after images (matlab)



### Theory / Process

First step is I used matlab (see the next page for code) to convert the given image to a text file. Next, to write the cuda converter, I imported the text file generated from matlab. I create a separate matrix for each rgb value and also create a grayscale matrix for final output. The data in the text file is stored by pixel, so that's why I have created matrices upfront. Another approach could have sent the whole matrix to the device, and then determined the pixels from there. I'm sure there would be a difference in performance between the two implementations. I then populate the rgb matrices and the allocate space for them on the device using `cudamalloc`, as well as the gray matrix. Then, I copy the rgb values to the device and set the dim3 number of blocks. Here I just use 1024 to keep things simple. Then, I call the kernel making sure to include the number of threads per block and the number of blocks in the call. After this the device will do its work and I then copy the matrix back onto the host. Lastly, I output the final grayscale matrix and then put it in matlab to check. See my matlab and cuda code on the next page.

## Matlab Code

```
% Image to text
rgb = imread('goats.jpeg');
writematrix(rgb, 'goats.txt', 'delimiter', ' ');
% Text to image
filename = 'grayscale.txt'
grayImage = uint8(importdata(filename));
imwrite(grayImage, 'myimage.png');
```

## Cuda Code

```
/*
    Nolan Anderson
    CPE 613 Week 2 Assignment 2
    This program takes in an image that has been converted to a text
    file.
    This text file has the RGB values. I converted the file in matlab
    using the following code:
    rgb = imread('goats.jpeg');
    writematrix(rgb, 'goats.txt', 'delimiter', ' ');
    ^^^^^ MAKE SURE TO PUT THE DEMINSIONS ON THE FIRST ROW OF THE TEXT
    FILE ^^^^^
    -----
    After this program, you should be able to run the following in
    matlab and see the grayscale image:
    filename = 'graygoats.txt'
    grayImage = uint8(importdata(filename));
    imwrite(grayImage, 'myimage.png');
    -----
    To compile, first 'load module cuda'
    Then: 'nvcc graygoats.cu -o goats'

    Next, 'run_gpu goats' and use all default values.
    This should result in a grayscale data of your original image.
*/

#include <fstream>
#include <iostream>
using namespace std;
```

```

// RGB to gray kernel. Pretty much just what was in the slides.
__global__ void rgb2gray_kernel(unsigned char* r, unsigned char* g,
unsigned char* b, unsigned char* gray, int matrixHeight, int
matrixWidth){
    unsigned int row = blockIdx.y*blockDim.y + threadIdx.y;    // Row
and column values
    unsigned int col = blockIdx.x*blockDim.x + threadIdx.x;
    unsigned int index = row*matrixWidth + col;
    if(row < matrixHeight && col < matrixWidth)                // Make
sure we don't go out of bounds on the row & column
        gray[index] = r[index] *3/10 + g[index]*6/10 + b[index]*1/10; //
General calculation for RGB -> gray values.
}

int main(){
    // Allocate matrices, first using matrixWidth and matrixHeight.
    // The data is coming from goats.txt, or any image whose data is in
    // text format.
    ifstream inFile;                                // input file
    int matrixWidth, matrixHeight;                  // matrixWidth and
matrixHeight of the matrix.
    inFile.open("goats.txt");                        // open the input file
    inFile >> matrixWidth >> matrixHeight; // pull the matrix
deminsions out of the file. (on the first line)
    const int matrixSize = matrixWidth * matrixHeight; // Obtain the
matrix size.
    unsigned char *r, *g, *b, *gray;                // Create the pointers to
hold the rgb data and the grayscale data.

    // Allocate a matrix for each color.
    unsigned char grayMatrix[matrixSize], rMatrix[matrixSize],
gMatrix[matrixSize], bMatrix[matrixSize];

    // Populate each matrix with the input file values
    int temp;
    for(int r = 0; r < matrixSize; r++){            // Loop through number of
values in matrix
        inFile >> temp;
        rMatrix[r] = temp;                          // Assign to each matrix.

```

```

        inFile >> temp;
        gMatrix[r] = temp;
        inFile >> temp;
        bMatrix[r] = temp;
    }

    // Allocate space on the device for rgb and gray matrices.
    // The size is the size of an unsigned char, multiplied by the size
of the matrix.
    cudaMalloc((void**)&r, matrixSize*sizeof(unsigned char));
    cudaMalloc((void**)&g, matrixSize*sizeof(unsigned char));
    cudaMalloc((void**)&b, matrixSize*sizeof(unsigned char));
    cudaMalloc((void**)&gray, matrixSize*sizeof(unsigned char));

    // Copy the rgb values to the device. Make sure to leave out the
gray value. This one will only be
    // copied back to the host.
    cudaMemcpy(r, &rMatrix, matrixSize * sizeof(unsigned char),
cudaMemcpyHostToDevice);
    cudaMemcpy(g, &gMatrix, matrixSize * sizeof(unsigned char),
cudaMemcpyHostToDevice);
    cudaMemcpy(b, &bMatrix, matrixSize * sizeof(unsigned char),
cudaMemcpyHostToDevice);

    // Create the dim3. 1024x1024 threads per block. Just using the
standard number of blocks from
    // the slides.
    dim3 numThreadsPerBlock(1024,1024);
    dim3 numBlocks((matrixWidth + numThreadsPerBlock.x -
1)/numThreadsPerBlock.x,
                    (matrixHeight + numThreadsPerBlock.y -
1)/numThreadsPerBlock.y);

    // Call the kernel, passing the appropriate values in. We need to
add the dim3 in the <<< >>> so the
    // warps / kernel / SMs know how to allocate their data. Gray is the
output and matrixheight/width are
    // there to support the calculations and indexing.
    rgb2gray_kernel <<< numThreadsPerBlock, numBlocks >>> (r, g, b,

```

```

gray, matrixHeight, matrixWidth);

    // Copy the data back to host from the device.
    cudaMemcpy(grayMatrix, gray, matrixSize*sizeof(unsigned char),
cudaMemcpyDeviceToHost);

    // Write the data to an outfile. In this case we're going back to a
text file.
    ofstream outFile;                // ofstream variable
    outFile.open("graygoats.txt"); // open the file
    for(int row = 0; row < matrixHeight; row++){           // Loop through
the rows
        for(int col = 0; col < matrixWidth; col++)         // Loop through
the columns
            outFile << +grayMatrix[row*matrixWidth + col] << " "; //
go to index to get current column in row.
        outFile << endl;    // Move to a new line (we're done with the
previous row)
    }

    // Close the files and exit program.
    inFile.close();
    outFile.close();
    return 0;
}

```