

Nolan Anderson | Week 2 Assignment 1 | CPE 613 | January 21 2023

- Implement the RGB to grayscale kernel and run it on a test JPEG image. Use MATLAB or Python functions to generate the RGB data from an image and view the before and after results. Upload a brief PDF report describing your process and showing your results.
- Use this image: [goats.jpeg](#)
- In Matlab, load it via `rgb = imread('goats.jpeg')` to get the RGB in a 456 x 810 x 3 uint8.
- Run `imshow(rgb)` to convince yourself this is an image of goats.
- Write this out to a file and load it into your program.
- Copy it up to the device.
- Convert to grayscale on the device.
- Copy the result down.
- Write the result to a file.
- Load the file into Matlab.
- Run `imshow` on the loaded data to see the grayscale image.

Before and after images (matlab)



Theory / Process

First, I used matlab (see the next page for code) to convert the given image to a text file. I then import the text file into the cuda program and create 4 matrix pointers, one for each rgb value and an empty gray matrix. These will be the destination matrices in the `cudamemcpy` call. To get the source for `cudamemcpy`, I create 4 additional empty matrices that are the size of the original matrix. I populate these matrices with the data in the text file. Next, I allocate space on the device using the first set of matrices with size `matrixsize`. Using these two sets, I copy the data to the device. Next, I create the `dim3` thread and block count. There are 1024 threads per block. Afterwards I call the kernel. This uses the individual rgb value matrices to do the calculation instead of having one big matrix with all the values. I then copy the data back to the host and output to the text file. See the matlab code below on how I converted the text file back into an image.

Matlab Code

```

% Image to text
rgb = imread('goats.jpeg');
writematrix(rgb,'goats.txt', 'delimiter', ' ');
% Text to image
filename = 'grayscale.txt'
grayImage = uint8(importdata(filename));
imwrite(grayImage, 'myimage.png');

```

Cuda Code

```

/*
    Nolan Anderson
    CPE 613 Week 2 Assignment 2
    This program takes in an image that has been converted to a text file.
    This text file has the RGB values. I converted the file in matlab using the following code:
    rgb = imread('goats.jpeg');
    writematrix(rgb,'goats.txt', 'delimiter', ' ');
    ^^^^^ MAKE SURE TO PUT THE DEMINSIONS ON THE FIRST ROW OF THE TEXT FILE ^^^^^
    -----
    After this program, you should be able to run the following in matlab and see the grayscale
    image:
    filename = 'graygoats.txt'
    grayImage = uint8(importdata(filename));
    imwrite(grayImage, 'myimage.png');
    -----
    To compile, first 'load module cuda'
    Then: 'nvcc graygoats.cu -o goats'

    Next, 'run_gpu goats' and use all default values.
    This should result in a grayscale data of your original image.
*/

#include <fstream>
#include <iostream>
using namespace std;

// RGB to gray kernel. Pretty much just what was in the slides.
__global__ void rgb2gray_kernel(unsigned char* r, unsigned char* g, unsigned char* b, unsigned char*
gray, int matrixHeight, int matrixWidth){
    unsigned int row = blockIdx.y*blockDim.y + threadIdx.y;    // Row and column values
    unsigned int col = blockIdx.x*blockDim.x + threadIdx.x;
    unsigned int index = row*matrixWidth + col;
    if(row < matrixHeight && col < matrixWidth)    // Make sure we don't go out of bounds
on the row & column
        gray[index] = r[index] *3/10 + g[index]*6/10 + b[index]*1/10; // General calculation for RGB
-> gray values.
}

int main(){
    // Allocate matrices, first using matrixWidth and matrixHeight.
    // The data is coming from goats.txt, or any image whose data is in
    // text format.
    ifstream inFile;    // input file
    int matrixWidth, matrixHeight;    // matrixWidth and matrixHeight of the matrix.

```

```

inFile.open("goats.txt"); // open the input file
inFile >> matrixWidth >> matrixHeight; // pull the matrix deminsions out of the file. (on the
first line)
const int matrixSize = matrixWidth * matrixHeight; // Obtain the matrix size.
unsigned char *r, *g, *b, *gray; // Create the pointers to hold the rgb data and the
grayscale data.

// Allocate a matrix for each color.
unsigned char grayMatrix[matrixSize], rMatrix[matrixSize], gMatrix[matrixSize],
bMatrix[matrixSize];

// Populate each matrix with the input file values
int temp;
for(int r = 0; r < matrixSize; r++){ // Loop through number of values in matrix
    inFile >> temp;
    rMatrix[r] = temp; // Assign to each matrix.
    inFile >> temp;
    gMatrix[r] = temp;
    inFile >> temp;
    bMatrix[r] = temp;
}

// Allocate space on the device for rgb and gray matrices.
// The size is the size of an unsigned char, multiplied by the size of the matrix.
cudaMalloc((void**)&r, matrixSize*sizeof(unsigned char));
cudaMalloc((void**)&g, matrixSize*sizeof(unsigned char));
cudaMalloc((void**)&b, matrixSize*sizeof(unsigned char));
cudaMalloc((void**)&gray, matrixSize*sizeof(unsigned char));

// Copy the rgb values to the device. Make sure to leave out the gray value. This one will only
be
// copied back to the host.
cudaMemcpy(r, &rMatrix, matrixSize * sizeof(unsigned char), cudaMemcpyHostToDevice);
cudaMemcpy(g, &gMatrix, matrixSize * sizeof(unsigned char), cudaMemcpyHostToDevice);
cudaMemcpy(b, &bMatrix, matrixSize * sizeof(unsigned char), cudaMemcpyHostToDevice);

// Create the dim3. 1024x1024 threads per block. Just using the standard number of blocks from
// the slides.
dim3 numThreadsPerBlock(1024,1024);
dim3 numBlocks((matrixWidth + numThreadsPerBlock.x - 1)/numThreadsPerBlock.x,
               (matrixHeight + numThreadsPerBlock.y - 1)/numThreadsPerBlock.y);

// Call the kernel, passing the appropriate values in. We need to add the dim3 in the <<< >>> so
the
// warps / kernel / SMs know how to allocate their data. Gray is the output and
matrixheight/width are
// there to support the calculations and indexing.
rgb2gray_kernel <<< numThreadsPerBlock, numBlocks >>> (r, g, b, gray, matrixHeight, matrixWidth);

// Copy the data back to host from the device.
cudaMemcpy(grayMatrix, gray, matrixSize*sizeof(unsigned char), cudaMemcpyDeviceToHost);

// Write the data to an outfile. In this case we're going back to a text file.
ofstream outFile; // ofstream variable
outFile.open("graygoats.txt"); // open the file

```

```
for(int row = 0; row < matrixHeight; row++){           // Loop through the rows
    for(int col = 0; col < matrixWidth; col++)          // Loop through the columns
        outFile << +grayMatrix[row*matrixWidth + col] << " "; // go to index to get current
column in row.
        outFile << endl;    // Move to a new line (we're done with the previous row)
    }

// Close the files and exit program.
inFile.close();
outFile.close();
return 0;
}
```