

# CPE348: Introduction to Computer Networks

## Lecture #17: Chapter 6

---



Jianqing Liu  
Assistant Professor of Electrical and Computer  
Engineering, University of Alabama in Huntsville

[jianqing.liu@uah.edu](mailto:jianqing.liu@uah.edu)  
<http://jianqingliu.net>

# TCP Congestion Control Overview

- Early on, we talked about:
  - TCP sliding window **protocol**;
  - TCP adaptive transmission **algorithms**.
- In this lecture, we will talk about:
  - TCP congestion control **protocols**
    - Additive Increase/Multiplicative Decrease
    - Slow Start
    - Fast Retransmit and Fast Recovery

# TCP Congestion Control – history

- It was introduced into the Internet in the late 1980s by Van Jacobson, ~ 8 yrs after the TCP/IP had become operational.
- Immediately preceding this time, the Internet was suffering from congestion collapse —
  - hosts would send their packets into the Internet as fast as the **advertised window** would allow
  - congestion would occur at some router, causing pkts drop
  - hosts would time out and retransmit their packets, resulting in even more congestion


# TCP Congestion Control – history

## ■ Basic idea

- Source determines how much capacity is available in the network,
  - it uses the arrival of an ACK as a signal that one of its packets has left the network,
  - source can then insert a new packet into the network.
- By using ACKs to pace the transmission of packets, TCP is said to be *self-clocking*.


# TCP Congestion Control

- Three popular TCP congestion control protocols
  - Additive Increase/Multiplicative Decrease
  - Slow Start
  - Fast Retransmit and Fast Recovery



# Additive Increase Multiplicative Decrease

- Source maintains a new variable **CongestionWindow**
- TCP's effective window is revised as follows:
  - **MaxWindow** = **MIN(CongestionWindow, AdvertisedWindow)**
  - **EffectiveWindow** =  
**MaxWindow – (LastByteSent – LastByteAcked).**
- That is, **MaxWindow** replaces **AdvertisedWindow**;
- Source is allowed to send no faster than the slowest component—the network or the destination host—can accommodate.



# Additive Increase Multiplicative Decrease

- Question is how source comes to learn an appropriate value for **CongestionWindow**.
  - First, source sets it as a default value based its perception to the network.
  - Then,
    - Decrease the **CongestionWindow** when the level of congestion goes up and
    - Increase the **CongestionWindow** when the level of congestion goes down.
    - Taken together, the mechanism is commonly called *additive increase/multiplicative decrease (AIMD)*

# Additive Increase Multiplicative Decrease

- The next question is how source comes to learn congestion occurs.
  - Observation: when packets are not delivered, timeout occurs
  - Then, source interprets timeouts as a sign of congestion and reduces the rate at which it is transmitting.
    - When a timeout occurs, the source sets **CongestionWindow** to **half** of its previous value – **multiplicative decrease**
    - When no timeout occurs, the source increments **CongestionWindow** by **1** – **additive increase**.




# Additive Increase Multiplicative Decrease

- For example
  - **CongestionWindow** is defined in terms of bytes;
  - Then,
    - Firstly, **CongestionWindow** is currently set to 16 packets. If a loss is detected, **CongestionWindow** is set to 8 packets.
    - Additional losses cause **CongestionWindow** to be reduced to 4, then 2, and finally to 1 packet.
    - **CongestionWindow** is not allowed to fall below the size of a single packet, or in TCP terminology, the *maximum segment size (MSS)*.



# Additive Increase Multiplicative Decrease

## ■ Another example



Packets in transit during additive increase, with one packet being added each RTT.


# Slow Start

- Drawback of AIMD: **too conservative!** It takes too long to ramp up a connection starting from scratch.
- TCP therefore provides a second mechanism, ironically called *slow start*,
  - Increase the **CongestionWindow** rapidly from a cold start when the **CongestionWindow** size starts at 1;
  - Slow start effectively increases the congestion window exponentially, rather than linearly (additive increase).

# Slow Start

- For example
  - Source starts out by setting **CongestionWindow** to 1 packet.
  - When the ACK for this packet arrives, source adds 1 to **CongestionWindow** and then sends 2 packets.
  - Upon receiving the corresponding 2 ACKs, source increments **CongestionWindow** by 2 and next sends 4 packets.
  - The result is that source doubles the number of packets it has in transit every RTT.

# Slow Start

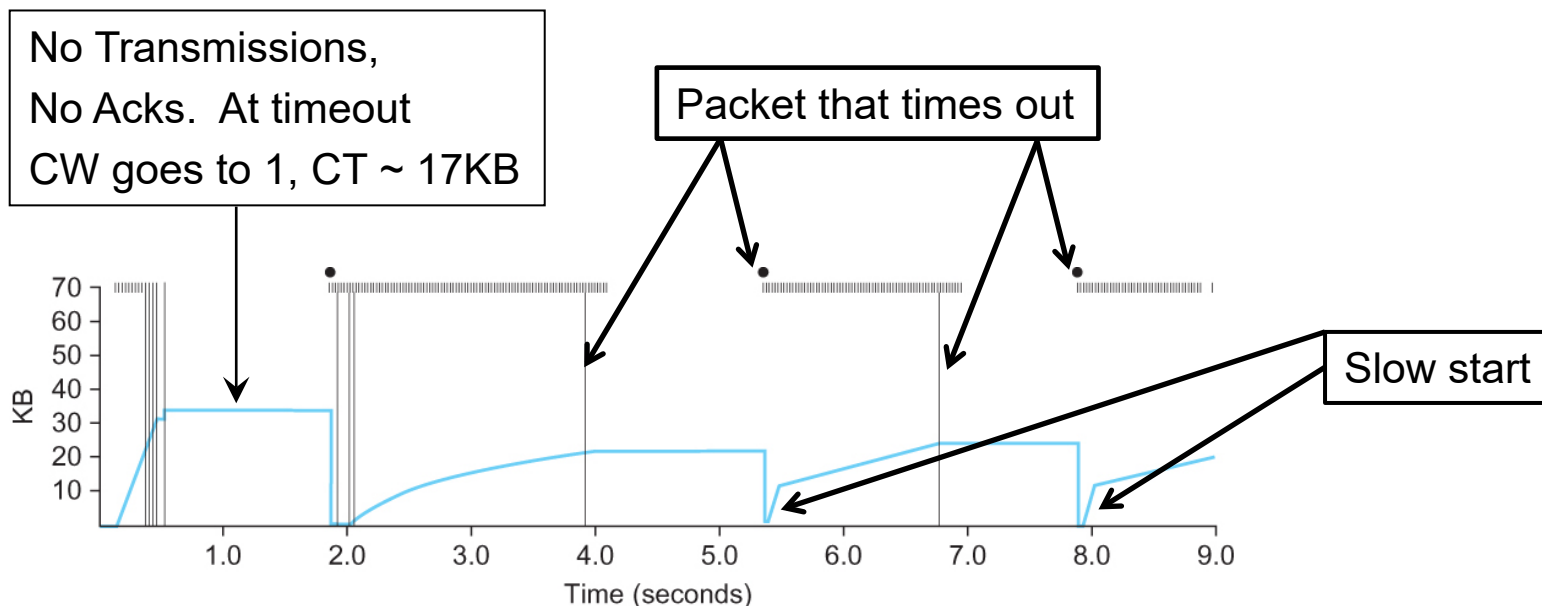


# Slow Start

- There are two situations in which slow start is **favorable**:
  - The first is at the very beginning of a connection:
    - Slow start continues to double **CongestionWindow** each RTT
    - After a packet is lost, a timeout causes multiplicative decrease to divide **CongestionWindow** by 2.
  - The second happens when the connection goes dead while waiting for a timeout to occur:
    - After communications for a while, source has a useful value of **CongestionWindow**, which is called **CongestionThreshold**;
    - Use **CongestionThreshold** as the “target” **CongestionWindow**.
    - Slow start is used to rapidly increase the sending rate up to this value, and then additive increase is used beyond this point. **(Combination! Fine-resolution!)**

# Slow Start


## ■ Cont' – Example for timeout occurs



Behavior of TCP congestion control. Colored line = value of CongestionWindow over time; solid bullets at top of graph = timeouts; hash marks at top of graph = time when each packet is transmitted; vertical bars = time when a packet that was eventually retransmitted was first transmitted.

# Fast Retransmit and Fast Recovery

- Issue of **slow start + AIMD**: TCP timeouts led to long periods of time during which the connection went dead while waiting for a timer to expire.



- A new mechanism called **fast retransmit** was added to TCP.
- It is a heuristic that sometimes triggers the retransmission of a dropped packet sooner than the timeout.



# Fast Retransmit and Fast Recovery

## ■ Basic idea

- Receiver responds with an ACK every time a packet arrives - even if it has already been ACKed
- When an out of order packet is received, TCP resends the same ACK as last time
- TCP cannot ACK the data contained in an out of order packet

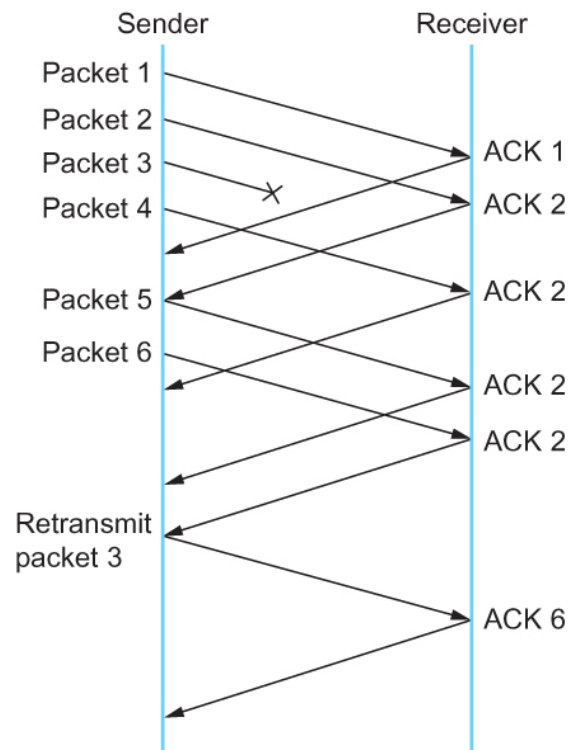
## Receiver Assisted Approach

# Fast Retransmit and Fast Recovery

- Basic idea – cont'
  - When the source sees a duplicate ACK,
    - It knows that the other side must have received a packet out of order,
    - Suggests that an earlier packet might have been lost.
  - In practice, TCP waits until it has seen three duplicate ACKs before retransmitting the packet.
    - in case the packet is delayed instead of lost


# Fast Retransmit and Fast Recovery

## ■ Example

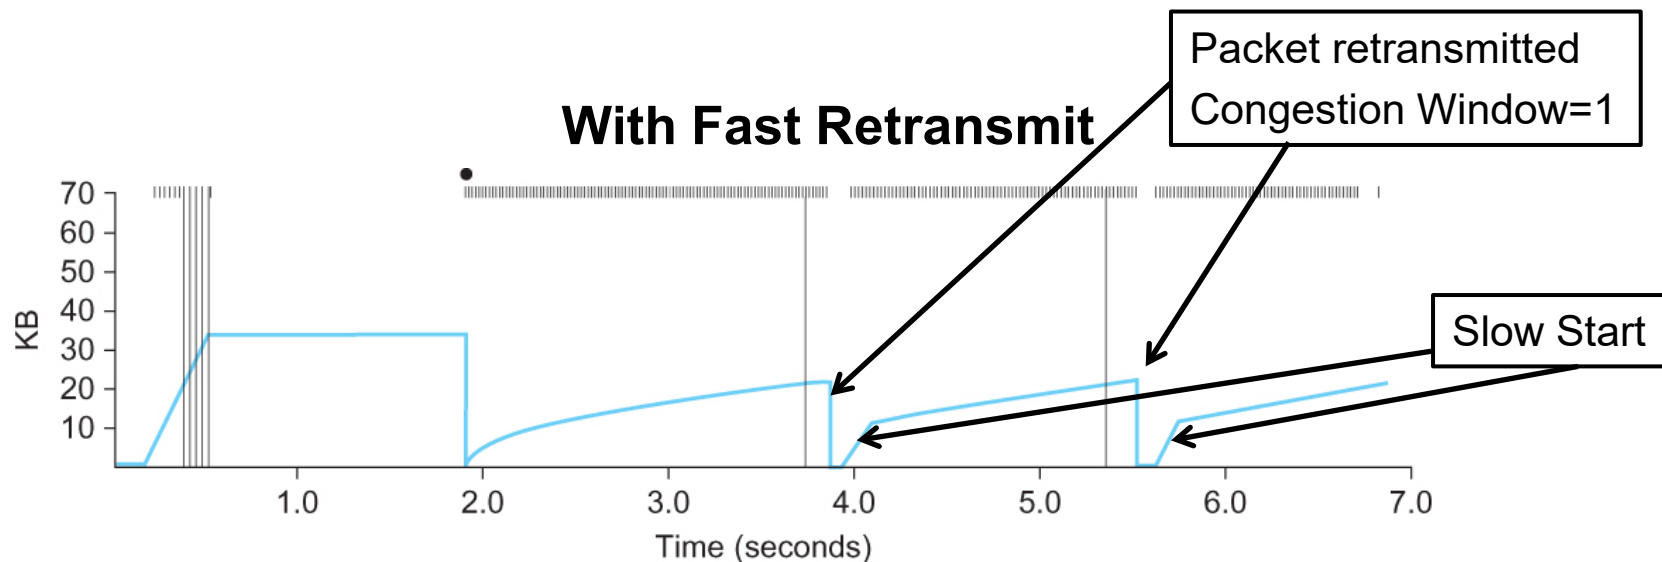


# Fast Retransmit and Fast Recovery

## Without Fast Retransmit



## With Fast Retransmit



# Fast Retransmit and Fast Recovery

## ■ Fast Recovery

- When the fast retransmit mechanism signals congestion,
  - NOT drop the **CongestionWindow** back to 1 packet and run slow start,
  - it is possible to use the ACKs that are still in the pipe to clock the sending of packets.
- This mechanism, which is called **fast recovery**.

# TCP Congestion Control Summary

- In this lecture, we will talk about:
  - TCP congestion control protocols;
    - Additive Increase/Multiplicative Decrease
    - Slow Start
    - Fast Retransmit and Fast Recovery
  - The (dis-)advantages of them;
  - You are expected to draw the tx diagram given a specific TCP congestion control protocol;