

CPE348: Introduction to Computer Networks

Lecture #18: Chapter 6

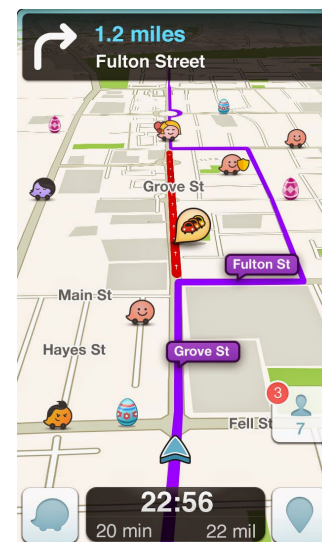


Jianqing Liu
Assistant Professor of Electrical and Computer
Engineering, University of Alabama in Huntsville

jianqing.liu@uah.edu
<http://jianqingliu.net>

TCP Congestion Control Overview

- Early on, we talked about:
 - TCP congestion control **protocols**;
- In this lecture, we will talk about:
 - TCP congestion avoidance **protocols**
 - DEC bit
 - Slow Start
 - Fast Retransmit and Fast Recovery



Congestion Avoidance Mechanism

- TCP congestion control is the strategy once congestion happens, as opposed to avoiding congestion in the first place.
- An appealing alternative is **congestion avoidance**
 - Has not yet been widely adopted,
 - Tries to predict when congestion is about to happen
 - Reduces the rate at which source sends data just before packets start getting lost.

Congestion Avoidance Mechanism

■ DEC Bit

- Initially developed for Digital Network Architecture (DNA), a connectionless network with a connection-oriented transport protocol;
- This mechanism could, also be applied to TCP/IP.
- The idea is to more evenly split the responsibility for congestion control between the **routers** and the **end nodes**.


Congestion Avoidance Mechanism

■ DEC Bit

- Each router monitors the load it is experiencing and explicitly notifies the end nodes when congestion is about to occur.
- This notification is implemented by setting a binary congestion bit in the packets that flow through the router; hence the name DECbit.
- The destination host then copies this congestion bit into the ACK it sends back to the source.
- Finally, the source adjusts its sending rate so as to avoid congestion

Congestion Avoidance Mechanism

- DEC Bit – router participation
 - The router calculates the average queue length over the specified interval.



Computing average queue length at a router

Congestion Avoidance Mechanism

- DEC Bit – source participation
 - The source records how many of its packets resulted in some router setting the congestion bit.
 - In particular, the source maintains a congestion window, just as in TCP
 - The source watches to see what fraction of the last window's worth of packets resulted in the congestion bit being set.

Congestion Avoidance Mechanism

- DEC Bit – source participation
 - If less than 50% of the packets had the bit set, then the source increases its congestion window by one packet.
 - If 50% or more of the last window's worth of packets had the congestion bit set, then the source decreases its congestion window to 0.875 times the previous value.

Congestion Avoidance Mechanism

- DEC Bit – source participation
 - 50% was chosen as the threshold based on analysis that showed it to correspond to the peak of the power curve.
 - The “increase by 1, decrease by 0.875” rule was selected because AIMD makes the mechanism stable.

Congestion Avoidance Mechanism

- Random Early Detection (RED)
 - A second congestion avoidance mechanism is called **random early detection (RED)**
 - When a router detects that congestion is imminent, it notifies the source to adjust its congestion window.

Congestion Avoidance Mechanism

- The difference between RED and DECbit
 - RED does not explicitly send a congestion notification message to the source
 - RED is implemented such that it *implicitly notifies* the source of congestion by dropping one of its packets.
 - The source is then notified by the subsequent timeout or duplicate ACK.

Congestion Avoidance Mechanism

■ Cont'

- The router drops the packet earlier than it has to
- This action notifies the source that it should decrease its congestion window sooner than without RED.
- The router drops a few packets before it has exhausted its buffer space completely
 - Causes the source to slow down its transmission rate
 - The router hopes that this will not cause to drop more packets later on.

Aggressive method!

Congestion Avoidance Mechanism

- Cont' – how to drop packets
 - Consider a simple FIFO queue on a router;
 - With RED, a router could decide to drop each arriving packet with some **drop probability** whenever the queue length exceeds some **drop level**.
 - This idea is called **early random drop**.

Congestion Avoidance Mechanism

- Cont' – Early Random Drop
 - RED computes an average queue length
 - **AvgLen** is iteratively computed as
$$\text{AvgLen} = (1 - \text{Weight}) \times \text{AvgLen} + \text{Weight} \times \text{SampleLen}$$
 - where $0 < \text{Weight} < 1$ and SampleLen is the length of the queue when a sample measurement is made.
 - Weight is a constant that determines how sensitive **AvgLen** is to variations in SampleLen
 - In most cases, the queue length is measured every time a new packet arrives at the router.


Congestion Avoidance Mechanism

■ Cont' – Early Random Drop

- Next, RED has two queue length thresholds : **MinThreshold** and **MaxThreshold**.
- When a packet arrives at the router, RED compares the current **AvgLen** with them, according to the following rules:
 - if $\text{AvgLen} \leq \text{MinThreshold}$
 - queue the packet
 - if $\text{MinThreshold} < \text{AvgLen} < \text{MaxThreshold}$
 - calculate probability P
 - drop the arriving packet with probability P
 - if $\text{AvgLen} \geq \text{MaxThreshold}$
 - drop the arriving packet

Congestion Avoidance Mechanism


- Cont' – Early Random Drop



RED thresholds on a FIFO queue

Congestion Avoidance Mechanism

- RED – how dropping probability is calculated:



Drop probability function for RED

Congestion Avoidance Mechanism

- Source-based Congestion Avoidance
 - Sources watch for some indication that congestion will happen soon;
 - For example, source may notice an increase in the RTT for each successive packet being sent.

Catch a sign!



Congestion Avoidance Mechanism

- Source-based Congestion Avoidance
 - One algorithm using RTT measurements:
 - The congestion window normally increases as in TCP
 - After every two round-trip delays, the algorithm
 - Averages the minimum and maximum RTT seen so far
 - Checks the current RTT against the average value
 - If the current RTT is greater than the average, source decreases the congestion window by $1/8$.

Congestion Avoidance Mechanism

- Source-based Congestion Avoidance
 - Another algorithm decides **CongestionWindow** based on changes to the RTT and the window size.
 - The congestion window is adjusted once every two RTTs based on the product
 - $(\text{CurrentWindow} - \text{OldWindow}) \times (\text{CurrentRTT} - \text{OldRTT})$
 - For a positive value, decrease the window size by 1/8;
 - For a non-positive value, increase the window by one maximum packet size.

Congestion Avoidance Mechanism

- Source-based Congestion Avoidance
 - A third algorithm the **CongestionWindow** based on the flatness of the sending rate
 - The **CongestionWindow** is adjusted every RTT based on the measured throughput
 - Every RTT increase the congestion window by 1
 - Compare the throughput achieved with the larger window to the throughput achieved with previous congestion window size.
 - If the difference is less than $\frac{1}{2}$ of the throughput when only one packet was in transit (measured at beginning of connection with congestion window at 1) then decrease window by 1.